

# Terraform files explanation

## Terraform files and explanation

The first five files have been pre-created from the gen-backend.sh script in the tf-setup stage, The S3 bucket and DynamoDB tables were also pre-created in the tf-setup stage.

### backend-cluster.tf, vars-main.tf

As described in previous sections.

---

### data-eks-cluster.tf

Get a data resource ("read only") reference for the EKS cluster control plane. Note the use of **data.terraform\_remote\_state.cluster.xxx** variables.

```
data "aws_eks_cluster" "eks_cluster" {
  name = data.terraform_remote_state.cluster.outputs.cluster-name
}

output "endpoint" {
  value = data.aws_eks_cluster.eks_cluster.endpoint
}

output "ca" {
  value = data.aws_eks_cluster.eks_cluster.certificate_authority[0].data
}

# Only available on Kubernetes version 1.13 and 1.14 clusters created or upgraded on or after September 3,
# 2019.
output "identity-oidc-issuer" {
  value = data.aws_eks_cluster.eks_cluster.identity[0].oidc[0].issuer
}
```

```
output "cluster-name" {
  value = data.aws_eks_cluster.eks_cluster.name
}
```

## user\_data.tf

This file will be base64 encoded and passed into the launch template is will:

- ❑ Join this node to the cluster **sudo /etc/eks/bootstrap.sh**
  - Note how some parameters for this are passed via Terraform data resources eg. `'${data.aws_eks_cluster.eks_cluster.name}'`
- ❑ Install our custom software/configuration - in this case the SSM agent.

```
locals {
  eks-node-private-userdata = <<USERDATA
  MIME-Version: 1.0
  Content-Type: multipart/mixed; boundary=="MYBOUNDARY=="

  ---MYBOUNDARY---
  Content-Type: text/x-shellscript; charset="us-ascii"

  #!/bin/bash -xe

  sudo /etc/eks/bootstrap.sh --apiserver-endpoint '${data.aws_eks_cluster.eks_cluster.endpoint}' --b64-
  cluster-ca '${data.aws_eks_cluster.eks_cluster.certificate_authority[0].data}'
  '${data.aws_eks_cluster.eks_cluster.name}'

  echo "Running custom user data script" > /tmp/me.txt

  yum install -y amazon-ssm-agent

  echo "yum'd agent" >> /tmp/me.txt

  systemctl enable amazon-ssm-agent && systemctl start amazon-ssm-agent

  date >> /tmp/me.txt

  ---MYBOUNDARY---
  USERDATA
}
```

## ssm-param-ami.tf

This gets the latest Amazon Linux 2 AMI for EKS from Systems Manager parameter store.

```
data "aws_ssm_parameter" "eksami" {
  name=format("/aws/service/eks/optimized-ami/%s/amazon-linux-2/recommended/image_id",
data.aws_eks_cluster.eks_cluster.version)
}
```

## launch\_template.tf

The launch template to use with the EKS managed node, this references:

- ❑ Our choice of AMI: **image\_id = data.aws\_ssm\_parameter.eksami.value.**
- ❑ Our base64 user data script **user\_data = base64encode(local.eks-node-private-userdata).**

The use of **create\_before\_destroy=true** is also important to allow us to create new versions of the launch template.

#### User data for worker launch

```
resource "aws_launch_template" "lt-ng1" {
  depends_on = [null_resource.auth_cluster]
  instance_type      = "t3.small"
  key_name           = "eksworkshop"
  name               = format("at-lt-%s-ng1", data.aws_eks_cluster.eks_cluster.name)
  tags               = {}
  image_id           = data.aws_ssm_parameter.eksami.value
  user_data          = base64encode(local.eks-node-private-userdata)
  vpc_security_group_ids = [data.terraform_remote_state.net.outputs.allnodes-sg]
  tag_specifications {
    resource_type = "instance"
  }
  tags = {
    Name = format("%s-ng1", data.aws_eks_cluster.eks_cluster.name)
  }
}
```

```

    }
    lifecycle {
      create_before_destroy=true
    }
  }
}

```

## aws\_eks\_node\_group\_\_manamieksp\_ng1.tf

```

resource "aws_eks_node_group" "ng1" {
  #ami_type    = "AL2_x86_64"
  depends_on  = [aws_launch_template.lt-ng1]
  cluster_name = data.aws_eks_cluster.eks_cluster.name
  disk_size   = 0
  instance_types = []
  labels = {
    "alpha.eksctl.io/cluster-name" = data.aws_eks_cluster.eks_cluster.name
    "alpha.eksctl.io/nodegroup-name" = format("ng1-%s", data.aws_eks_cluster.eks_cluster.name)
  }
  node_group_name = format("ng1-%s", data.aws_eks_cluster.eks_cluster.name)
  node_role_arn   = data.terraform_remote_state.iam.outputs.nodegroup_role_arn
  #release_version = "1.17.11-20201007"
  subnet_ids = [
    data.terraform_remote_state.net.outputs.sub-priv1,
    data.terraform_remote_state.net.outputs.sub-priv2,
    data.terraform_remote_state.net.outputs.sub-priv3,
  ]
  tags = {
    "alpha.eksctl.io/cluster-name"      = data.aws_eks_cluster.eks_cluster.name
    "alpha.eksctl.io/eksctl-version"    = "0.29.2"
    "alpha.eksctl.io/nodegroup-name"    = format("ng1-%s", data.aws_eks_cluster.eks_cluster.name)
    "alpha.eksctl.io/nodegroup-type"    = "managed"
    "eksctl.cluster.k8s.io/v1alpha1/cluster-name" = data.aws_eks_cluster.eks_cluster.name
  }
  #version = "1.17"
}

```

```

launch_template {
  name   = aws_launch_template.lt-ng1.name
  version = "1"
}

scaling_config {
  desired_size = 2
  max_size     = 3
  min_size     = 1
}

lifecycle {
  ignore_changes = [scaling_config[0].desired_size]
}

timeouts {}
}

```

## **null\_resource.tf**

The null resource runs the test.sh and auth.sh script after the creation of the cluster **depends\_on** = [aws\_eks\_cluster.cluster]

```

resource "null_resource" "gen_cluster_auth" {
  triggers = {
    always_run = timestamp()
  }
  depends_on = [aws_eks_node_group.ng1]
  provisioner "local-exec" {
    on_failure = fail
    when       = create
    interpreter = ["/bin/bash", "-c"]
    command    = <<EOT
      ./c9-auth.sh
      ./auth-cicd.sh
    >>EOT
  }
}

```

```

    echo "*****"
EOT
}
}

```

## c9-auth.sh

Authorize the local user to the cluster via ~/.kube/config

```

test -n "$C9_PID" && echo C9_PID is "$C9_PID" || "echo C9_PID is not set && exit"
echo "local auth"
sleep 5
c9builder=$(aws cloud9 describe-environment-memberships --environment-id=$C9_PID | jq -r
'memberships[].userArn')
if echo ${c9builder} | grep -q user; then
    rolearn=${c9builder}
    echo Role ARN: ${rolearn}
elif echo ${c9builder} | grep -q assumed-role; then
    assumedrolename=$(echo ${c9builder} | awk -F/ '{print $(NF-1)}')
    rolearn=$(aws iam get-role --role-name ${assumedrolename} --query Role.Arn --output text)
    echo Role ARN: ${rolearn}
fi
## need to Terraform this ?
cat << EOF > patch.yaml
data:
  mapUsers: |
    - userarn: ${rolearn}
      username: admin
  groups:
    - system:masters
EOF
kubectl get configmap -n kube-system aws-auth -o yaml > aws-auth.yaml
cat patch.yaml >> aws-auth.yaml
kubectl apply -f aws-auth.yaml

```

```
#eksctl create iamidentitymapping --cluster mycluster1 --arn ${rolearn} --group system:masters --username admin
```

## auth-cicd.sh

This script authorizes the CodeBuild role to access the EKS cluster by patching the aws-auth configmap

The CodeBuild role **codebuild-eks-cicd-build-app-service-role** is added to the `system:masters` kubernetes group which gives full admin rights to the cluster.

In production environments you would want to scope this down to perhaps a specific namespace using Kubernetes RBAC.

```
test -n "$ACCOUNT_ID" && echo ACCOUNT_ID is "$ACCOUNT_ID" || "echo ACCOUNT_ID is not set && exit"
ROLE="  - rolearn: arn:aws:iam::$ACCOUNT_ID:role/codebuild-eks-cicd-build-app-service-role\nusername: build\n  groups:\n    - system:masters"
#
kubectl get -n kube-system configmap/aws-auth -o yaml | awk "/mapRoles: \\/\\/{print;print\n\"$ROLE\\\";next}1" > /tmp/aws-auth-patch.yml
#
kubectl patch configmap/aws-auth -n kube-system --patch "$(cat /tmp/aws-auth-patch.yml)"
```

## output.tf

Some output variables are defined, but they are not used in this workshop

Expand to view Terraform code

```
locals {
  config-map-aws-auth = <<CONFIGMAPAWSAUTH
  apiVersion: v1
  kind: ConfigMap
  metadata:
    name: aws-auth
    namespace: kube-system
  data:
```

```
mapRoles: |
- rolearn: data.terraform_remote_state.iam.outputs.nodegroup_role_arn
  username: system:node:{{EC2PrivateDNSName}}
  groups:
    - system:bootstrappers
    - system:nodes
```

CONFIGMAPAWSAUTH

```
kubeconfig = <<KUBECONFIG
apiVersion: v1
clusters:
- cluster:
    server: aws_eks_cluster.eks-cluster.endpoint
    certificate-authority-data: aws_eks_cluster.eks-cluster.certificate_authority.0.data
  name: kubernetes
contexts:
- context:
    cluster: kubernetes
    user: aws
  name: aws
current-context: aws
kind: Config
preferences: {}
users:
- name: aws
  user:
    exec:
      apiVersion: client.authentication.k8s.io/v1alpha1
      command: aws-iam-authenticator
      args:
        - "token"
        - "-i"
        - "aws_eks_cluster.eks-cluster.name"
KUBECONFIG
}
```



```
output "config-map-aws-auth" {  
  value = "local.config-map-aws-auth"  
}
```

```
output "kubeconfig" {  
  value = "local.kubeconfig"  
}
```