

Terraform files explanation

Terraform files and explanation

The first three files have been pre-created from the gen-backend.sh script in the tf-setup stage.

backend-sample.tf & vars-main.tf

As previously described.

k8s.tf

Include the kubernetes provider

```
provider "kubernetes" {}
```

Define the namespace using Terraform, notice the extended timeout period for deletion **delete = "20m"** to allow ingress resources time to delete.

sampleapp-namespace.tf

```
resource "kubernetes_namespace" "game-2048" {  
  metadata {  
    name = "game-2048"  
  }  
  
  timeouts {  
    delete = "20m"  
  }  
}
```

sampleapp-deployment.tf

Define the deployment. There are a few notable parts to this file:

Two data resources are specified that return the current AWS region and account id

```
data "aws_region" "current" {} data "aws_caller_identity" "current" {}
```

These data resources are then used to construct a string for the image location.

```
image = format("%s.dkr.ecr.%s.amazonaws.com/sample-app",  
data.aws_caller_identity.current.account_id, data.aws_region.current.name)
```

There's also a node selector to ensure the app runs on our specific node group

```
node_selector = { "alpha.eksctl.io/nodegroup-name" = "ng1-mycluster1" }
```

```
data "aws_region" "current" {}
```

```
data "aws_caller_identity" "current" {}
```

```
resource "kubernetes_deployment" "game-2048_deployment-2048" {
```

```
  metadata {  
    name      = "deployment-2048"  
    namespace = "game-2048"  
  }  
}
```

```
  spec {  
    replicas = 4  
    selector {  
      match_labels = {  
        "app.kubernetes.io/name" = "app-2048"  
      }  
    }  
    strategy {  
      type = "RollingUpdate"  
    }  
  }  
}
```

```

rolling_update {
  max_surge    = "25%"
  max_unavailable = "25%"
}
}

template {
  metadata {
    annotations = {}
    labels     = { "app.kubernetes.io/name" = "app-2048" }
  }

  spec {

    node_selector      = { "alpha.eksctl.io/nodegroup-name" = "ng1-mycluster1" }
    restart_policy     = "Always"
    share_process_namespace = false
    termination_grace_period_seconds = 30

    container {
      image      = format("%s.dkr.ecr.%s.amazonaws.com/sample-app",
data.aws_caller_identity.current.account_id, data.aws_region.current.name)
      image_pull_policy = "Always"
      name             = "app-2048"
      port {
        container_port = 80
        host_port      = 0
        protocol       = "TCP"
      }

      resources {
      }
    }
  }
}
}
}

```

sampleapp-service.tf

Define the service:

```
resource "kubernetes_service" "game-2048__service-2048" {

  metadata {
    name      = "service-2048"
    namespace = "game-2048"
  }

  spec {
    selector = {
      "app.kubernetes.io/name" = "app-2048"
    }

    type = "NodePort"

    port {
      port      = 80
      protocol  = "TCP"
      target_port = "80"
    }
  }
}
```

sampleapp-ingress.tf

Define the ingress resource - this creates a AWS Application Load Balancer via the previously installed aws-load-balancer-controller.

Note how multiple annotations are passed, including the listener port 8080.

```
resource "kubernetes_ingress" "game-2048_ingress-2048" {
  metadata {
    annotations = {"kubernetes.io/ingress.class" = "alb", "alb.ingress.kubernetes.io/scheme" = "internal",
"alb.ingress.kubernetes.io/target-type" = "ip", "alb.ingress.kubernetes.io/listen-ports" = "[{"HTTP": 8080}]}"}
    name      = "ingress-2048"
    namespace = "game-2048"
  }

  spec {

    rule {
      http {
        path {
          path = "/"
          backend {
            service_name = "service-2048"
            service_port = "80"
          }
        }
      }
    }
  }
}
```
